

# exercise\_05

December 11, 2023

## 1 Exercise 5

```
[2]: import numpy as np
import matplotlib.pyplot as plt
```

### 1.1 Statistical testing: Theory

- 1) In statistical hypothesis testing, what does a *significant difference* mean? Use an example for illustration, e.g., the difference in weight of domestic and feral cats.
- 2) Given the same example as you used in 1), what is your null hypothesis ( $H_0$ ) and what is your alternate hypothesis ( $H_1$ )?
- 3) When conducting a statistical test, one important outcome is a so called *p-value*. What is a p-value? Explain using the same example as you used in 1).
- 4) Now, we have compared the weight of domestic and feral cats using a statistical test: what does a p-value of  $p=0.01$  mean?
- 5) What level ( $\alpha$ ) of p-value would you use if you did a statistical test that causes ...
  - 1) ... someone to go to jail.
  - 2) ... a patient to take cough sirup.
- 6) Going back to our cat example,
  - 1) if the calculated p-value does not pass your  $\alpha$  level, does that mean that the weight of domestic and feral cats is NOT different? Explain.
  - 2) if the calculated p-value passes your  $\alpha$  level, what do you conclude?
- 7) What is a Type I and a Type II error?

### 1.2 Statistical tests: Examples

```
[3]: from scipy.stats import ranksums, ttest_ind
```

#### 1.2.1 Comparing the difference of two (normal) distributions

- 1) What is the null hypothesis of the Wilcoxon Rank-Sum test (also known as Mann-Whitney U test)?


- 2) The following code generates data from two distributions and calculates the  $p$ -value based on the Wilcoxon Rank-Sum test. Try different `n_samples`, and `shift`. Describe your observations.
- 3) Could you also use the t-test for this data? Check the corresponding [Wikipedia entry, Section 'Uses'](#).

```
[4]: # settings
n_samples = 100
shift = 0.2
variance = 1

# generate data
v1 = np.random.normal(size=n_samples)
v2 = np.random.normal(shift, variance, size=n_samples)
plt.hist(v1, alpha=0.5)
plt.hist(v2, alpha=0.5);

# perform the test
ranksums(v1, v2)
```

```
[4]: RanksumsResult(statistic=0.19058433319583615, pvalue=0.8488512655032039)
```



### 1.2.2 Correlations

- 1) Intuitively, what does the Pearson correlation coefficient measure? What is its minimum and maximum value and what do different values mean?
- 2) How to interpret a p-value calculated for Pearson correlation?
- 3) In what way is Spearman correlation different from Pearson?
- 4) Look at the generated data below and run the code. Try different `n_samples` and `covariance` values. What do you observe?
- 5) Add the outlier of (0, 99) to the data. What do you observe?
- 6) Calculate the Spearman correlation (`scipy.stats.spearmanr`) as well as Kendall's tau (`scipy.stats.kendalltau`). How does the outlier influence these correlation measures?

```
[5]: from scipy.stats import spearmanr, pearsonr, kendalltau
```

```
[6]: # try different sample sizes and covariances

n_samples = 50


# use values from -1 to 1
covariance = 0

# generate data
d = np.random.multivariate_normal((1,1), [[1,covariance],[covariance,1]],  
    ↪n_samples)
v1 = d[:,0]
v2 = d[:,1]

# plot
fig, axes = plt.subplots(1,2, figsize=(2 * 5, 1 * 5))
ax = axes[0]; ax.hist2d(v1, v2)
ax = axes[1]; ax.scatter(v1, v2)

# stats
r, p = pearsonr(v1, v2)
print(f"r={r}, p={p}")
```

```
r=0.06992005432325617, p=0.6294538874318618
```



### 1.3 Multiple hypothesis comparison

- 1) Explain what the multiple hypothesis comparison is and why it is a problem (refer to p-values in the process).
- 2) Execute the following code, experiment with `n_samples` and `p_threshold`, and explain what it does.

```
[8]: # multiple hypothesis comparison and correction on real data
# try different p_thresholds (also called alpha values) and sample sizes

# parameters
n_samples = 1000
n_features = 1000
p_threshold = 0.05 # alpha value

# generate data
X = np.random.random((n_samples, n_features))
y = np.random.random(n_samples)

# calculate p-values between all features and the target
ps = []
for i in range(X.shape[1]):
    r, p = spearmanr(X[:,i], y)
    ps.append(p)

# report results
```

```
print(f"# of tests with p < {p_threshold} without correction: ", (np.array(ps) <
    ↪ p_threshold).sum())

from statsmodels.stats.multitest import multipletests
ps2 = multipletests(ps, method="fdr_bh")[1]
print(f"# of tests with p < {p_threshold} WITH correction: ", (np.array(ps2) <
    ↪ p_threshold).sum())
```

# of tests with p < 0.05 without correction: 46

# of tests with p < 0.05 WITH correction: 0

## 1.4 Model comparison

- 1) Run the given code. What does it do?
- 2) Adjust the given code to use *repeated* cross-validation (`StratifiedRepeatedKFold`) for three different models (add one model). Why are we using a `stratified` version here?
- 3) Run the training several times with different numbers of repetitions (1, 10, 100). What do you observe?
- 4) Compare your performance values using `Wilcoxon rank-sum test`. Again, compare different numbers of repetitions. What do you observe?
- 5) Is one of the two models significantly better? Why?
- 6) Retrain the “better” model on the whole training dataset and plot your predictions (`predict_proba`) for the training and test data.

```
[9]: # imports
import numpy as np
import pandas as pd

from sklearn.model_selection import train_test_split, cross_validate,
    ↪ StratifiedKFold, RepeatedStratifiedKFold

from sklearn.pipeline import make_pipeline

from sklearn.impute import SimpleImputer
from sklearn.preprocessing import StandardScaler

from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.svm import LinearSVC

import matplotlib.pyplot as plt
import seaborn as sns
```

```
[10]: # load data
data_titanic = pd.read_csv("exercise_01_intro-to-python_titanic.csv",
    ↪ index_col="PassengerId")
```

```
[11]: def extract_features(data):
        """Extract features from existing variables"""

        data_extract = data.copy()

        # name
        name_only = data_extract["Name"].str.replace(r"\(.*\)", "", regex=True)
        first_name = name_only.str.split(", ", expand=True).iloc[:,1]
        title = first_name.str.split(".", expand=True).iloc[:,0]
        data_extract["Title"] = title

        # ticket

        # ...

        return data_extract

data_extract = extract_features(data_titanic)

[12]: def preprocess(data):
        """Convert features into numeric variables readable by our models."""

        data_preprocessed = data.copy()

        # Sex

        data_preprocessed = pd.get_dummies(data_preprocessed, columns=["Sex"],
        ↪drop_first=True)

        # Embarked

        data_preprocessed = pd.get_dummies(data_preprocessed, columns=["Embarked"],
        ↪dummy_na=True)

        # Title
        title = data_preprocessed["Title"]
        title_counts = title.value_counts()
        higher_titles = title_counts[title_counts < 50]
        title_groups = ["higher" if t in higher_titles else t for t in title]
        data_preprocessed["Title"] = title_groups
        data_preprocessed = pd.get_dummies(data_preprocessed, columns=["Title"])

        # drop the rest
        data_preprocessed.drop(columns=["Name", "Cabin", "Ticket"], inplace=True)

        return data_preprocessed
```

```
data_preprocessed = preprocess(data_extract)
```

```
[13]: # before inspecting the data, selecting and building models, etc.  
# FIRST split data into train and test data (we set the test data size to 30%)  
  
X = data_preprocessed.drop(columns="Survived")  
y = data_preprocessed["Survived"]  
  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,  
↳stratify=y)
```

```
[14]: def evaluate(model, name, scoring="accuracy"):  
    results = cross_validate(  
        model,  
        X_train, y_train,  
        cv=StratifiedKFold(n_splits=5),  
        scoring=scoring,  
        return_train_score=True)  
  
    results = pd.DataFrame(results)  
    results["model"] = name  
  
    display(f"Train ({scoring}): {results['train_score'].mean():.02f} +/-  
↳{results['train_score'].std():.02f}")  
    display(f"Test ({scoring}): {results['test_score'].mean():.02f} +/-  
↳{results['test_score'].std():.02f}")  
    return results
```

```
[16]: results = []
```

```
[17]: imputer = SimpleImputer(strategy="median")  
scaler = StandardScaler()  
model_lr = LogisticRegression()  
  
pipeline_lr = make_pipeline(imputer, scaler, model_lr)  
results_lr = evaluate(pipeline_lr, "logistic regression")  
results.append(results_lr)
```

```
'Train (accuracy): 0.81 +/- 0.01'
```

```
'Test (accuracy): 0.80 +/- 0.03'
```

```
[18]: imputer = SimpleImputer(strategy="median")  
model_tree = DecisionTreeClassifier(max_depth=3)  
  
pipeline_tree = make_pipeline(imputer, model_tree)  
results_dt = evaluate(pipeline_tree, "decision tree")  
results.append(results_dt)
```

```
'Train (accuracy): 0.81 +/- 0.01'
```

```
'Test (accuracy): 0.80 +/- 0.03'
```

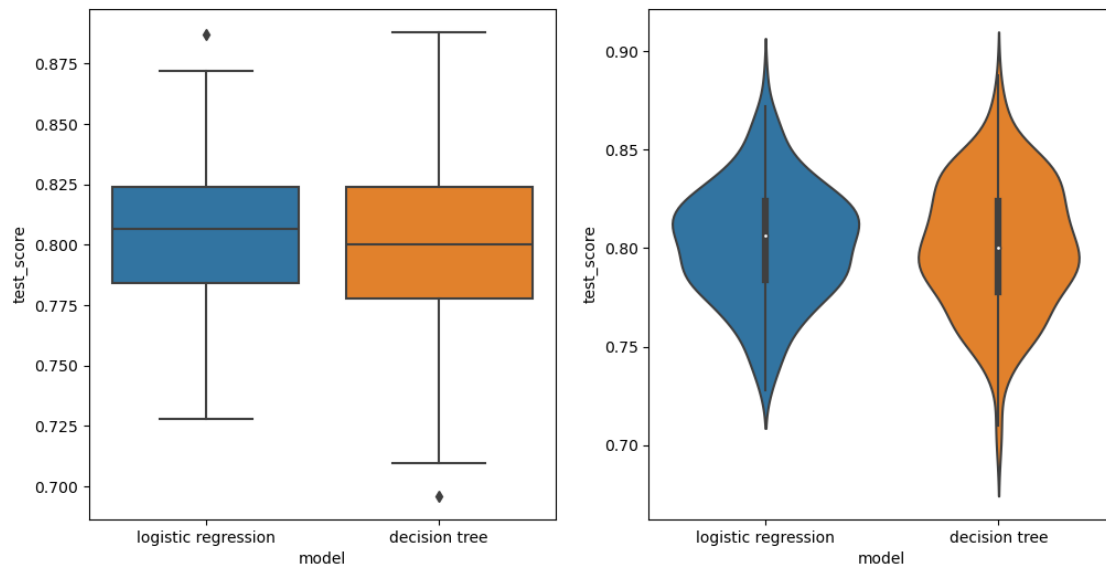
```
[19]: # # another model
# imputer = SimpleImputer(strategy="median")
# model_sum = SVC()

# pipeline_sum = make_pipeline(imputer, model_sum)
# results_sum = evaluate(pipeline_sum, "sum")
# results.append(results_sum)
```

```
[20]: results_df = pd.concat(results)

fig, axes = plt.subplots(1, 2, figsize=(2 * 6, 6))
ax = axes[0]
sns.boxplot(data=results_df, x="model", y="test_score", ax=ax)
ax = axes[1]
sns.violinplot(data=results_df, x="model", y="test_score", ax=ax)
```

```
[20]: <Axes: xlabel='model', ylabel='test_score'>
```



```
[21]: # compare
from scipy.stats import ranksums

test_score_lr = results_df.loc[results_df["model"] == "logistic regression",
    ↪ "test_score"]
test_score_dt = results_df.loc[results_df["model"] == "decision tree",
    ↪ "test_score"]
```



```
# ...
```

## 1.5 Model introspection

The following code plots feature importances for a linear model.


- 1) Try different data splits. What do you observe?
- 2) Try adding and removing individual features. What do you observe?
- 3) Try scaling individual features. What do you observe?
- 4) Are the features that are considered important according to a model the only possibly important features to make predictions?

```
[24]: # before inspecting the data, selecting and building models, etc.  
# FIRST split data into train and test data (we set the test data size to 30%)  
  
X = data_preprocessed.drop(columns="Survived")  
y = data_preprocessed["Survived"]  
  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,  
↪stratify=y)
```

```
[25]: imputer = SimpleImputer(strategy="mean")  
scaler = StandardScaler()  
preprocessing_pipeline = make_pipeline(scaler, imputer)  
  
X_train_example = X_train.copy()  
  
# TODO: drop features here  
# ...  
  
X_train_processed = preprocessing_pipeline.fit_transform(X_train_example)  
  
# TODO: change `X_train_processed` for example here  
# ...  
  
# set model  
model = LogisticRegression()  
  
# fit for coefficient extraction  
model.fit(X_train_processed, y_train)  
feature_weights = model.coef_.flatten()
```

```
[27]: # plot feature importances  
fig, ax = plt.subplots(1, 1, figsize=(6 * 2, 6), sharey=True)  
sns.barplot(x=X_train_example.columns, y=feature_weights, ax=ax)  
ax.set_xticklabels(ax.get_xticklabels(), rotation=270)
```

```
plt.show()
```



```
[ ]:
```