

exercise__06__v2

December 18, 2023

1 Exercise 6

1.1 Recall

Basics: List four typical methods applied in the context of Knowledge Discovery from Databases (KDD) applications and briefly describe them!

Basics: List and describe the 6 phases of the CRISP-DM model!

Basics: What should happen during the evaluation phase?

Evaluation: Define the accuracy measure.

Evaluation: What is class imbalance and why can it cause a high accuracy even though the classifier is bad? Give an example.

Evaluation: In what applications is precision more important than recall, and in which applications is recall more important than precision?

Evaluation: What is grouped cross validation and when should you use it? Come up with an example and iterate all splits.

Statistics: Multiple hypothesis comparison

- 1) Explain what the multiple hypothesis comparison is and why it is a problem (refer to p-values in the process).
- 2) What is a way to correct for multiple hypothesis comparison. Give a concrete example.
- 3) What is the difference between an α -value and a p-value?

1.2 Univariate analysis

1.2.1 Load the data

```
[2]: # imports
import numpy as np
import pandas as pd

from sklearn.model_selection import train_test_split, cross_validate, \
    StratifiedKFold, RepeatedStratifiedKFold

from sklearn.pipeline import make_pipeline

from sklearn.impute import SimpleImputer
```

```

from sklearn.preprocessing import StandardScaler

from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.svm import LinearSVC

import matplotlib.pyplot as plt
import seaborn as sns

```

```

[3]: # load data
data_titanic = pd.read_csv("exercise_01_intro-to-python_titanic.csv",
    ↪index_col="PassengerId")

```

```

[4]: def extract_features(data):
    """Extract features from existing variables"""

    data_extract = data.copy()

    # name
    name_only = data_extract["Name"].str.replace(r"\(.*\)", "", regex=True)
    first_name = name_only.str.split(", ", expand=True).iloc[:,1]
    title = first_name.str.split(".", expand=True).iloc[:,0]
    data_extract["Title"] = title

    # ticket

    # ...

    return data_extract

data_extract = extract_features(data_titanic)

```

```

[5]: def preprocess(data):
    """Convert features into numeric variables readable by our models."""

    data_preprocessed = data.copy()

    # Sex

    data_preprocessed = pd.get_dummies(data_preprocessed, columns=["Sex"],
    ↪drop_first=True)

    # Embarked

    data_preprocessed = pd.get_dummies(data_preprocessed, columns=["Embarked"],
    ↪dummy_na=True)

```

```

# Title
title = data_preprocessed["Title"]
title_counts = title.value_counts()
higher_titles = title_counts[title_counts < 50]
title_groups = ["higher" if t in higher_titles else t for t in title]
data_preprocessed["Title"] = title_groups
data_preprocessed = pd.get_dummies(data_preprocessed, columns=["Title"])

# drop the rest
data_preprocessed.drop(columns=["Name", "Cabin", "Ticket"], inplace=True)

return data_preprocessed

data_preprocessed = preprocess(data_extract)

```

```

[6]: # before inspecting the data, selecting and building models, etc.
# FIRST split data into train and test data (we set the test data size to 30%)

X = data_preprocessed.drop(columns="Survived")
y = data_preprocessed["Survived"]

```

```

[7]: X.head(3)

```

```

[7]:
      Pclass  Age  SibSp  Parch    Fare  Sex_male  Embarked_C \
PassengerId
1          3  22.0     1     0   7.2500      True     False
2          1  38.0     1     0  71.2833      False     True
3          3  26.0     0     0   7.9250      False     False

      Embarked_Q  Embarked_S  Embarked_nan  Title_Miss  Title_Mr \
PassengerId
1          False         True         False         False      True
2          False         False         False         False     False
3          False         True         False          True     False

      Title_Mrs  Title_higher
PassengerId
1          False         False
2           True         False
3          False         False

```

1.2.2 Univariate comparison

```

[8]: from scipy.stats import spearmanr, mannwhitneyu, ranksums

```

```
[9]: # for each feature calculate whether there is a statistical difference between
      ↪ survivors (1) and victims (0)
tests = {}
for c in X.columns:
    # TODO: fill this in by computing the test statistic for each feature
    ↪ columns!
    # tests[c] = mannwhitneyu()
    pass
tests
```

```
[9]: {}
```

Which is the feature most associated with surviving?

Is the association with Survived positive or negative?


```
[12]: # calculate the spearman correlation between all features
tests = {}
for c1 in X.columns:
    for c2 in X.columns:
        # TODO: fill this on
        # tests[(c1, c2)] = spearmanr()
        pass
tests
```

```
[12]: {}
```

```
[14]: # we can do the same with a pandas function:
corr = X.corr(method="spearman")
```

```
[15]: # and plot it
sns.heatmap(corr, center=0, vmin=-1, vmax=1)
```

```
[15]: <Axes: >
```



What is the maximum and minimum value of the Spearman correlation coefficient and what do they mean?

Answer: -1 (strong negative correlation) and 1 (strong positive correlation)

What is the strongest correlation of all feature pairs?

Answer: Check `tests`

1.3 BONUS:

Plot a scatter plot with each point representing a feature (e.g., using t-SNE for dimensionality reduction), the size representation the association to **Survived**, the color corresponding to the direction of the association, and lines between the features if their absolute correlation exceeds 0.5.

[]: